

Overview of Ledis - a lightweight Redis

Application: [here](#)

Choosing frontend

I quite unfamiliar with implementing a CLI, so first step I googled for some example. Then I discover [web-cli](#). Look at its features I though this is kind of overkill.

Then I remember working with a nice, simple Ruby console: [web-console](#), which is exactly what I need.

After I spend half an hour interate web-console into a Rails application. Web-console is not a front-end plugin I look into its source code, then pick the visual code that I need.

I have to read along the line and remove a bunch of thing Rails related like console session, key-word auto suggestion.

Choosing backend

I consider between Ruby and Elixir. Reason:

- Ruby: I familiar with it
- Elixir: Redis operation seem very functional-programming to me. And with the reputation of concurrency, I though Elixir would give greate performance to Ledis.

I spend nearly 1 day to make decision. But I realize, I don't have enough infomation to decide. To be honest, I'm lost!!!. Then I calm down and make a different approach: research more about Redis architecture

Here are interesting things I found about Redis:

1. Redis is single-threaded

before proper researching about redis, I was to implementing ledis with a multiprocess architecture in mind. The reason behind this is because I mostly use Ruby MRI. So to me, in order to scale throughput, I have to have multiple processes (along with some threading, but the point is I still have to use many processes).

The problem is each process has its own memory and is not shared to each other. So I wonder: how can Redis make sure that every requests have access to the same data storage? And this is when I realize, that Redis is actually single-threaded, and you can have concurrency on single-thread using an I/O multiplexing mechanism and an event loop

=> So I decided to go with Ruby to simplify the architecture

2. Every Redis command is atomic

Every command of Redis is atomic because of the fact that Redis is single-threaded. Ruby MRI have the GIL so that no more than 1 thread is run at a time. But this case is not equivalent to single-thread, because MRI actually have multi-thread, just no more than 1 thread run at a time -> must ensure single-thread. In order to do this, I config Puma server in single mode and have only 1 thread max

Ledis structure

Ledis class

This is the main class. Its objective is to hold a class variable *value_table*, which is hash. This hash will contain all keys and values

This hash is not shared with anything, so Ledis class also the place to implement the core logic.

Ledis::Command class

command_string inside request get parsed into *command* and *params*. A Ledis::Command object will be initialized with these information.

Objective of this class is

1. To check if this is a supported Ledis command.
2. Most of Ledis command are type-specific. This class will check if current command and key are compatible with each other.

3. Execute Ledis command. Those *command* and *params* will be passed to Ledis to perform command logic

Ledis::DataTypes

Please look at objective No.2 of Command. In order to know command-key capability, a naive way is get value of that key from the main datastore (the class variable hash) and know which type it is. Then look at a mapping between commands is the type it suppose to work with

I find this not very efficient, because:

- main datastore contain many info, so access key is slow
- have to invoke ruby code `object.class` each time.

Solution: create another class variable hash *type_table* which also have keys, but value is just in enum of data stype

With this solution, I was able to cut as much as half of the time for this command-key checking operation.

- Because *type_table* is lighter than *value_table*, It's only store in enum value
- Does not need to use Ruby to check type
- Downside: To cost more memory. But I think this is a good tradeoff, because I think we are going to read a lot more than write, and command-key checking operation is required for most of the command

If I had more time

There are lots of room for Ledis to be improved. If I had more time, I will:

- refactor front-end component to a more appropriate structure to allow static caching. Reason that I can not right now is web-console code is mainly a javascript file, with `js.erb` handler. Inside this file the author use Rails to render some html code, which is the UI component the CLI. An approach with BackboneJS view or ReactJS can archive the goal
- not using Rails at all. Because, again, this is overkill. a Rack application should be good enough. Right now I'm using Grape for API, which is also a Rack application.

Thankyou

Thank you for the opportunity. I've learned alot new things through this assignment, aside with strengthen my knowledge on concurrency.